

High-performance Computing, Autumn 2024

Robert Buj

November 5, 2024

Question 1

The graph below shows the execution time of the parallel implementation (OpenMP) of the sum of the different elements of the vector a ($sum = \sum_{n=0}^{N-1} a_i$) using the reduction clause. sum time means T_{PAR} and, global time means $T_{PAR} + T_{SEQ}$. To simplify the source code, the last two sentences that should be part of T_{PAR} are omitted.




Figure 1: Execution time for adding 1000000000 elements

```

#include <stdio.h>
#include <stdlib.h> /* rand, srand*/
#include <time.h>   /* clock */
#include <omp.h>     /* omp_get_wtime */
#define NANO_SECONDS 1000000000.0f

int main(int argc, const char *argv[]) {
    double start_sum_time, start_global_time, end_sum_time, end_global_time;
    float *a, sum = 0.0f;
    size_t i, n;

    start_global_time = omp_get_wtime();
    if (argc < 2) {
        printf("Usage: %s SIZE\n", argv[0]);
        exit(-1);
    }
    n = (size_t)abs(atoi(argv[1]));
    a = (float *)malloc(n * sizeof(float));
    srand((unsigned int)clock());
    for (i = 0; i < n; i++)
        /* random number -100,100 */
        a[i] = (((float)rand() / (float)RAND_MAX) - 0.5f) * 200.0f;
    start_sum_time = omp_get_wtime();
#ifdef USE_REDUCTION
    printf("using omp reduction...\n");
#pragma omp parallel for shared(a, n) private(i) reduction(+: sum)
#else
    printf("not using omp reduction...\n");
#endif
    for (i = 0; i < n; i++)
        sum += a[i];
    end_sum_time = omp_get_wtime();
    free(a);
    end_global_time = omp_get_wtime();
    /* size,sum_time,global_time,result */
    printf("%lu,%f,%f,%.2f\n", n,
        end_sum_time - start_sum_time,
        end_global_time - start_global_time, sum);
    exit(0);
}

```

Figure 2: sv.c

Question 2

In the figures below, you can see the execution times of the multiplication of two matrices. The Figure 3 summarizes the execution time for mm, and the Figure 4 summarizes the execution time for mm2.

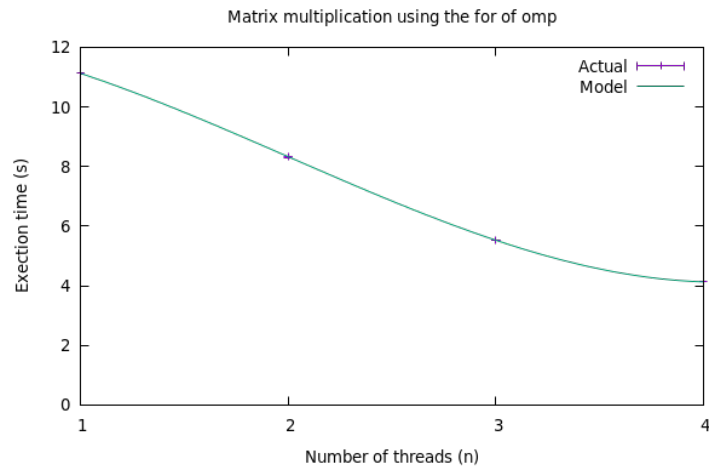


Figure 3: Execution time of mm

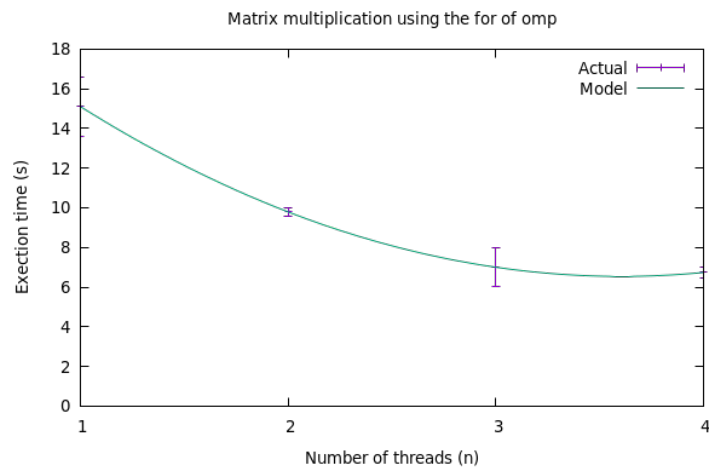


Figure 4: Execution time of mm2

```

#include <stdlib.h>
#include <stdio.h>
#ifdef _OPENMP
#include <omp.h>
#endif
#define SIZE 1000

int main()
{
    float matrixa[SIZE][SIZE], matrixb[SIZE][SIZE], mresult[SIZE][SIZE];
    int i, j, k;

    /* Initialize the Matrix arrays */
    for (i = 0; i < SIZE * SIZE; i++)
    {
        mresult[0][i] = 0.0f;
        matrixa[0][i] = matrixb[0][i] = ((float)rand()) * 1.1f;
    }

    /* Matrix-Matrix multiply */
#ifdef _OPENMP
#pragma omp parallel for private(i, j, k) shared(matrixa, matrixb, mresult)
#endif
    for (i = 0; i < SIZE; i++)
        for (j = 0; j < SIZE; j++)
            for (k = 0; k < SIZE; k++)
                mresult[i][j] = mresult[i][j] + matrixa[i][k] * matrixb[k][j];

    exit(0);
}

```

Figure 5: mm.c

```

#include <stdlib.h>
#include <stdio.h>
#ifdef _OPENMP
#include <omp.h>
#endif
#define SIZE 1000

int main()
{
    float matrixa[SIZE][SIZE], matrixb[SIZE][SIZE], mresult[SIZE][SIZE];
    int i, j, k;

    /* Initialize the Matrix arrays */
    for (i = 0; i < SIZE * SIZE; i++)
    {
        mresult[0][i] = 0.0f;
        matrixa[0][i] = matrixb[0][i] = ((float)rand()) * 1.1f;
    }

    /* Matrix-Matrix multiply */
#ifdef _OPENMP
#pragma omp parallel for private(i, j, k) shared(matrixa, matrixb, mresult)
#endif
    for (k = 0; k < SIZE; k++)
        for (j = 0; j < SIZE; j++)
            for (i = 0; i < SIZE; i++)
                mresult[i][j] = mresult[i][j] + matrixa[i][k] * matrixb[k][j];

    exit(0);
}

```

Figure 6: mm2.c

Question 3

In the second file there are more cache misses. On each iteration or very frequently, the data block that contains the element $mresult[i][j]$ has to be loaded from main memory to cache as well as $matrix[i][k]$, and then the result has to be stored from cache to memory ($mresult[i][j]$), and that has a time penalty. The deeper loop changes the value of i all time ($j*i*k$ times), it means that it changes the row of the matrices $mresult$ and $matrix$, which is not a consecutive memory position, thus it can't be found on existing stored data in cache and hence the cache miss.

	flops	flops2	factor
Real_time	11.158243	13.867895	1.24
Proc_time	11.1389094	13.8486805	1.24
Total_flops	2000768595	2003602758	=
MFLOPS	179.6227218	145.9744919	0.81

Table 1: Resume of flops.

host	Real_time	Proc_time	Total_flops	MFLOPS
compute-0-0.local	11.086303	11.060439	2000306880	180.852402
compute-0-1.local	11.144365	11.135186	2000525211	179.65799
compute-0-2.local	11.109171	11.083703	2000449652	180.485672
compute-0-2.local	11.177109	11.151365	2001101522	179.44902
compute-0-4.local	11.181869	11.1729	2001144338	179.106979
compute-0-5.local	11.173346	11.14662	2001155989	179.530289
compute-0-6.local	11.105769	11.080009	2000310428	180.533279
compute-0-7.local	11.191671	11.182329	2001156712	178.957062
compute-0-8.local	11.201117	11.192089	2000388215	178.73233
compute-0-9.local	11.21171	11.184454	2001146996	178.922195
Mean	11.158243	11.1389094	2000768595	179.6227218

Table 2: flops.csv

host	Real_time	Proc_time	Total_fpins	MFLOPS
compute-0-0.local	15.989087	15.952642	2005237273	125.699379
compute-0-1.local	12.904942	12.89483	2002888637	155.324936
compute-0-3.local	13.104773	13.094281	2003821725	153.030289
compute-0-3.local	15.979366	15.967595	2005076636	125.571609
compute-0-4.local	15.863821	15.851385	2005022654	126.488792
compute-0-5.local	12.886873	12.85668	2002693066	155.77063
compute-0-6.local	12.896558	12.86628	2002751814	155.658966
compute-0-7.local	12.867961	12.858011	2002615980	155.748505
compute-0-8.local	13.17238	13.162561	2003174048	152.187256
compute-0-9.local	13.867895	12.98254	2002745747	154.264557
Mean	13.867895	13.8486805	2003602758	145.9744919

Table 3: flops2.

Question 4

Mainly, any counter that allows you to see data cache misses such as PAPI_L2_DCM counter, or total cache misses such as PAPI_L3_TCM counter.

```
$ grep 'PAPI_L[1\|2\|3]_[D\|T]CM' papi_avail.1157538.out | grep Yes
PAPI_L1_DCM 0x80000000 Yes No Level 1 data cache misses
PAPI_L2_DCM 0x80000002 Yes Yes Level 2 data cache misses
PAPI_L1_TCM 0x80000006 Yes Yes Level 1 cache misses
PAPI_L2_TCM 0x80000007 Yes No Level 2 cache misses
PAPI_L3_TCM 0x80000008 Yes No Level 3 cache misses
```

Figure 7: Cache misses counters

Another approach would be to get the hit rate ($\text{hits} / (\text{misses} + \text{hits})$), for instance L2 data cache hit rate is equal to $\text{PAPI_L2_DCH} / (\text{PAPI_L2_DCM} + \text{PAPI_L2_DCH})$ or L2 total cache hit rate is equal to $\text{PAPI_L2_TCH} / (\text{PAPI_L2_TCM} + \text{PAPI_L2_TCH})$

```
$ grep hits papi_avail.1157538.out | grep Yes | grep -v instruction
PAPI_L2_DCH 0x8000003f Yes Yes Level 2 data cache hits
PAPI_L2_TCH 0x80000056 Yes Yes Level 2 total cache hits
$ grep 'PAPI_L2_DCM\|PAPI_L2_TCM' papi_avail.1157538.out
PAPI_L2_DCM 0x80000002 Yes Yes Level 2 data cache misses
PAPI_L2_TCM 0x80000007 Yes No Level 2 cache misses
```

Figure 8: Available hit rates

Question 5

Level 2 data cache misses

```
...
#define SIZE 1000
#define NUM_EVENTS 3

int main() {
    float matrixa[SIZE][SIZE], matrixb[SIZE][SIZE], mresult[SIZE][SIZE];
    int i,j,k, ret, events[NUM_EVENTS] = {
        PAPI_L2_DCM, /* Level 2 data cache misses */
        PAPI_TOT_INS,
        PAPI_FP_OPS
    };
    long long values[NUM_EVENTS];

    if (PAPI_num_counters() < NUM_EVENTS) {
        fprintf(stderr, "No hardware counters here, or PAPI not supported.\n");
        exit(1);
    }
    if ((ret = PAPI_start_counters(events, NUM_EVENTS)) != PAPI_OK) {
        fprintf(stderr, "PAPI failed to start counters: %s\n",
            PAPI_strerror(ret));
        exit(1);
    }
    ...
    if ((ret = PAPI_read_counters(values, NUM_EVENTS)) != PAPI_OK) {
        fprintf(stderr, "PAPI failed to read counters: %s\n",
            PAPI_strerror(ret));
        exit(1);
    }
    printf("Level 2 data cache misses = %lld\n", values[0]);
    printf("Total instructions %lld\n", values[1]);
    printf("Total hardware flops = %lld\n", values[2]);
    exit(0);
}
```

Figure 9: PAPI_L2_DCM counter

host	l2_data_cache_misses	flops	instructions
compute-0-3.local	73427877	2002301797	33073987835
compute-0-3.local	73468648	2002970061	33073987856
compute-0-4.local	73447817	2002914214	33073987295
compute-0-4.local	73475291	2002939833	33073987802
compute-0-5.local	73496434	2002970056	33073987844
compute-0-6.local	73516754	2002183809	33073987790
compute-0-6.local	73532795	2002190223	33073987271
compute-0-8.local	73446486	2002212906	33073987803
compute-0-8.local	73483198	2002231005	33073987315
compute-0-9.local	73403306	2002188586	33073987321
Mean	73469861	2002510249	33073987613

Table 4: counters.v1

host	l2_data_cache_misses	flops	instructions
compute-0-0.local	247343577	2007636067	33073992529
compute-0-1.local	198451499	2005369253	33073988880
compute-0-1.local	246925839	2004916107	33073989052
compute-0-2.local	217514011	2005463084	33073988940
compute-0-5.local	221126604	2007500382	33073992482
compute-0-6.local	282752381	2004950462	33073989134
compute-0-7.local	308986582	2004938915	33073989113
compute-0-9.local	207060775	2007283871	33073989206
compute-0-9.local	220119372	2005551652	33073989127
compute-0-9.local	242368607	2005456802	33073989249
Mean	239264925	2005906660	33073989771

Table 5: counters2.v1

	counters	counters2	factor
l2_data_cache_misses	73.469.861	239.264.925	3.25
flops	2.002.510.250	2.005.906.660	=
instructions	33.073.987.614	33.073.989.249	=

Level 3 cache misses

```
...
#define SIZE 1000
#define NUM_EVENTS 3

int main() {
...
    float matrixa[SIZE][SIZE], matrixb[SIZE][SIZE], mresult[SIZE][SIZE];
    int i,j,k,ret,events[NUM_EVENTS] = {
        PAPI_L3_TCM, /* Level 3 total cache misses */
        PAPI_TOT_INS,
        PAPI_FP_OPS
    };
    long long values[NUM_EVENTS];

    if (PAPI_num_counters() < NUM_EVENTS) {
        fprintf(stderr, "No hardware counters here, or PAPI not supported.\n");
        exit(1);
    }
    if ((ret = PAPI_start_counters(events, NUM_EVENTS)) != PAPI_OK) {
        fprintf(stderr, "PAPI failed to start counters: %s\n",
            PAPI_strerror(ret));
        exit(1);
    }
...
    if ((ret = PAPI_read_counters(values, NUM_EVENTS)) != PAPI_OK) {
        fprintf(stderr, "PAPI failed to read counters: %s\n",
            PAPI_strerror(ret));
        exit(1);
    }
    printf("Level 3 total cache misses = %lld\n", values[0]);
    printf("Total instructions %lld\n", values[1]);
    printf("Total hardware flops = %lld\n", values[2]);
    exit(0);
}
```

Figure 10: PAPI_L3_TCM counter

host	l3_cache_misses	flops	instructions
compute-0-0.local	5366017	2002226528	33073987289
compute-0-1.local	5678066	2003073334	33073987838
compute-0-2.local	5328895	2002995277	33073987313
compute-0-3.local	4260167	2002921687	33073987266
compute-0-4.local	5339415	2002248383	33073987290
compute-0-5.local	5338447	2002278852	33073987284
compute-0-7.local	5281872	2002917592	33073987310
compute-0-8.local	5174189	2003000856	33073987337
compute-0-8.local	5527467	2002254118	33073987332
compute-0-9.local	5765099	2002974628	33073987367
Mean	5305963	2002689126	33073987363

Table 6: counters.v2

host	l3_cache_misses	flops	instructions
compute-0-0.local	23552310	2006732564	33073992325
compute-0-1.local	15057050	2005263565	33073989022
compute-0-2.local	25044014	2007044985	33073992458
compute-0-3.local	14136625	2005319250	33073989008
compute-0-3.local	16348467	2006623896	33073988802
compute-0-5.local	17910968	2006887724	33073988795
compute-0-6.local	18309773	2006912908	33073989143
compute-0-6.local	21645323	2006613810	33073992559
compute-0-7.local	18873257	2005298686	33073989149
compute-0-8.local	24027023	2008809433	33073992592
Mean	19490481	2006550682	33073990385

Table 7: counters2.v2

	counters	counters2	factor
l3_total_cache_misses	5.305.964	19.490.481	3.67
flops	2.002.689.126	2.006.550.683	=
instructions	33.073.987.363	33.073.990.386	=

Question 6

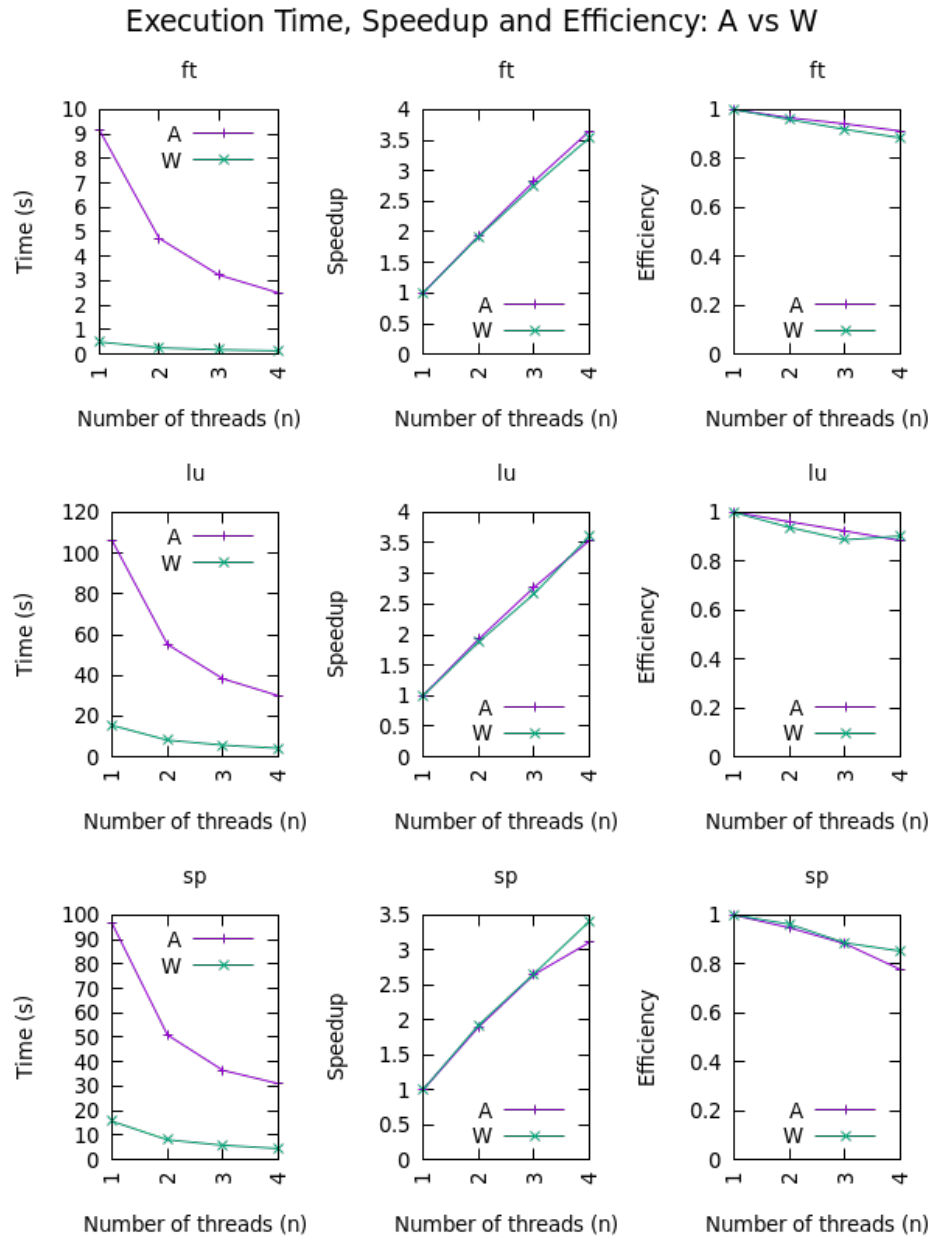


Figure 11: Problem class A and W

Execution Time, Speedup and Efficiency: A

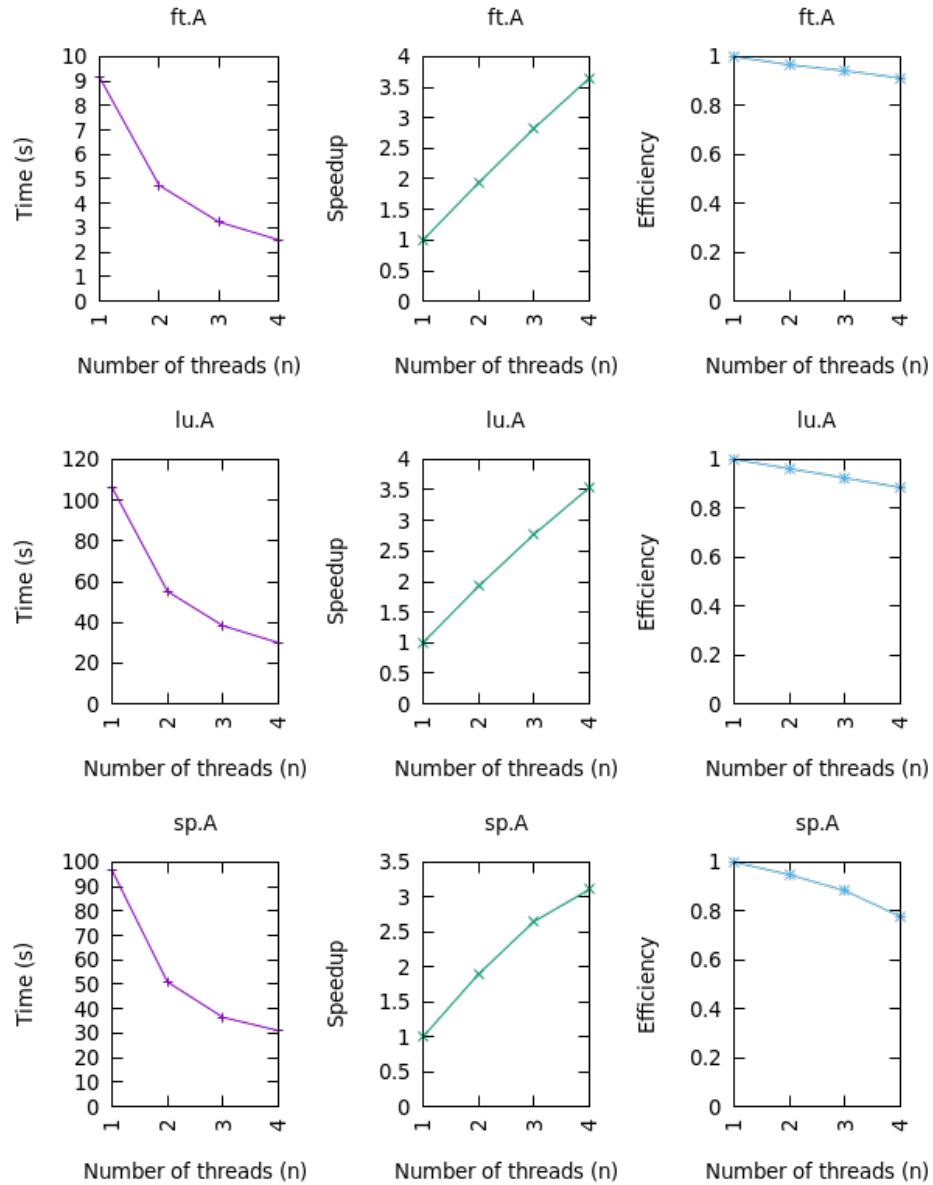


Figure 12: Problem class: A

Execution Time, Speedup and Efficiency: W

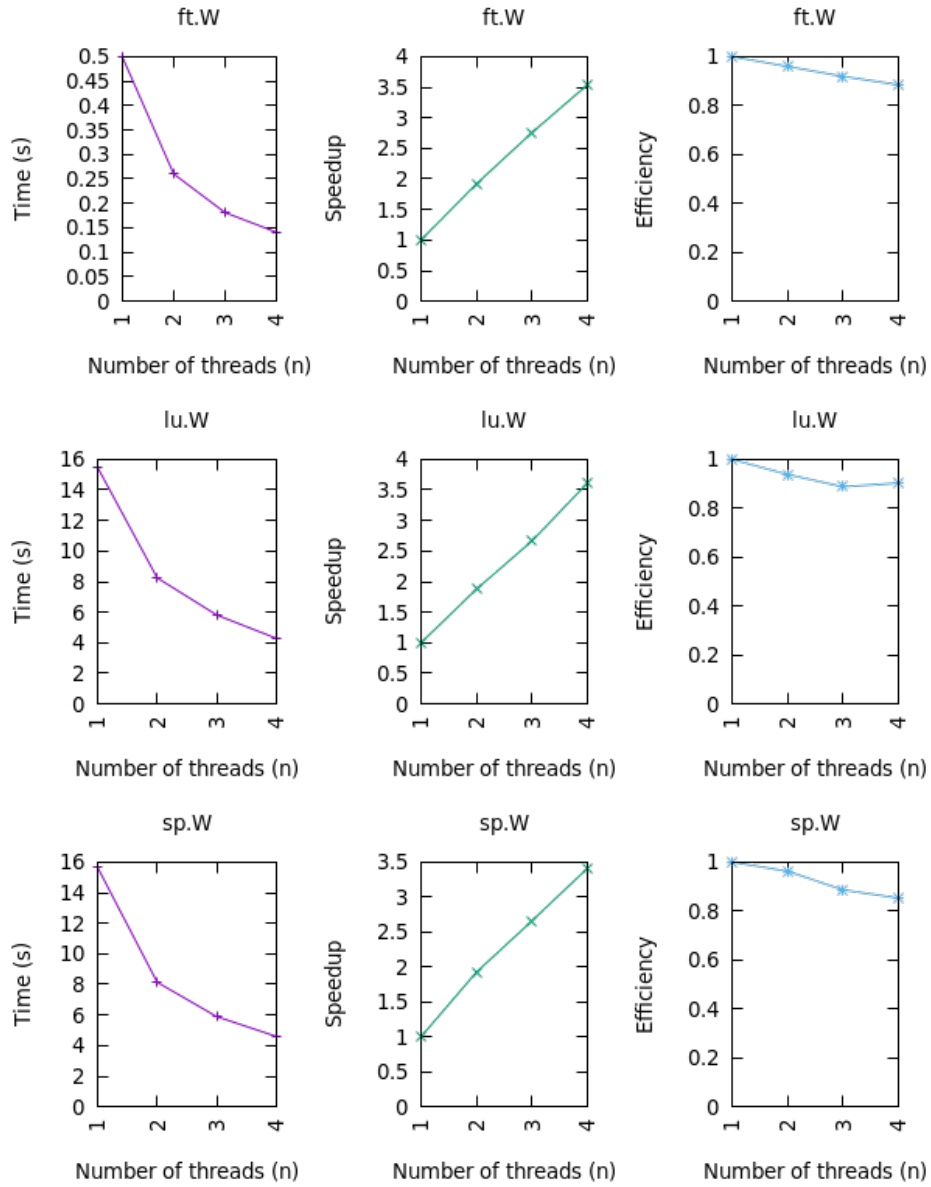


Figure 13: Problem class: W

In the execution time graph, a small value is better than a larger one. In the speed-up graph, a large value is better than a smaller one. In the efficiency graph, a value close to 1 is better than another closer to 0.

In all kernels, regardless of the problem size, it seems that the execution time graph is an exponential decrease function, the speed-up graph is a linear increase function with a high slope, and the efficiency graph is a slightly pronounced linear decrease function.

The speed-up and efficiency graphs are quite similar regarding the problem size, having worse results on sp kernel with problem size A.

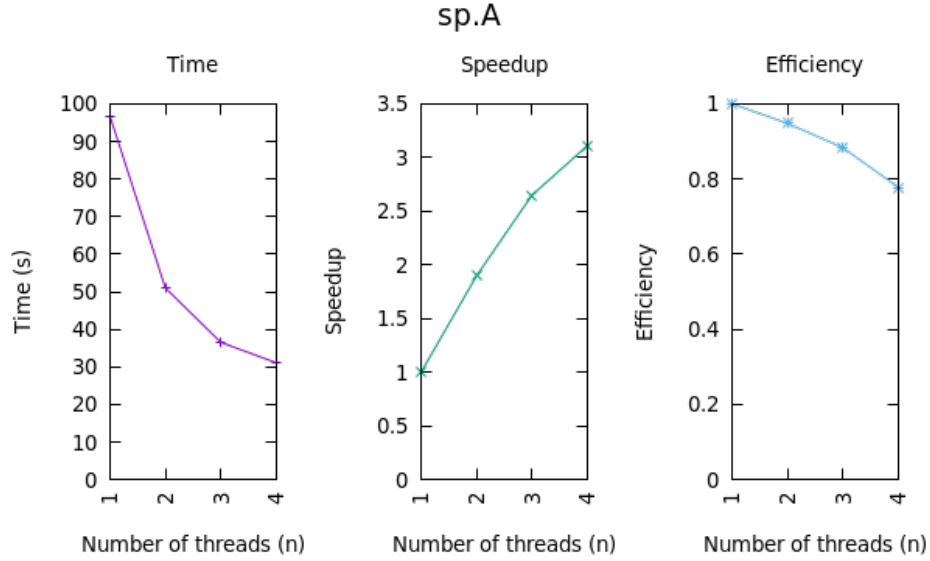


Figure 14: Kernel sp class A

On the other hand, it seems that the execution time graphs are quite similar between kernels, it means that sp and lu kernels have quite similar values for their execution times, and ft kernel has proportional ones.

Question 7

Scheduling is a method in OpenMP to distribute iterations to different threads in for loop.¹

The single-threaded execution time is the same on both kernels, about 15.5 s. However, the parallel execution with four threads has different improvement level according to the applied scheduling policy.

The sp kernel reaches higher speed-ups than lu. Almost scheduling policies applied to lu kernel reach about 3.5 of speed-up, on the other hand the lu kernel has an average speed-up value about 1,75 (excluding static scheduling), which is exactly half the speed-up value of sp kernel.

As the chunk-size increases, the speed-up decreases on both kernels. It seems that for a given chunk-size, the speed-up reaches the same level for dynamic and guided policies.




Figure 15: Execution Time lu and sp

¹Yiling, (2020), OpenMP - Scheduling(static, dynamic, guided, runtime, auto)

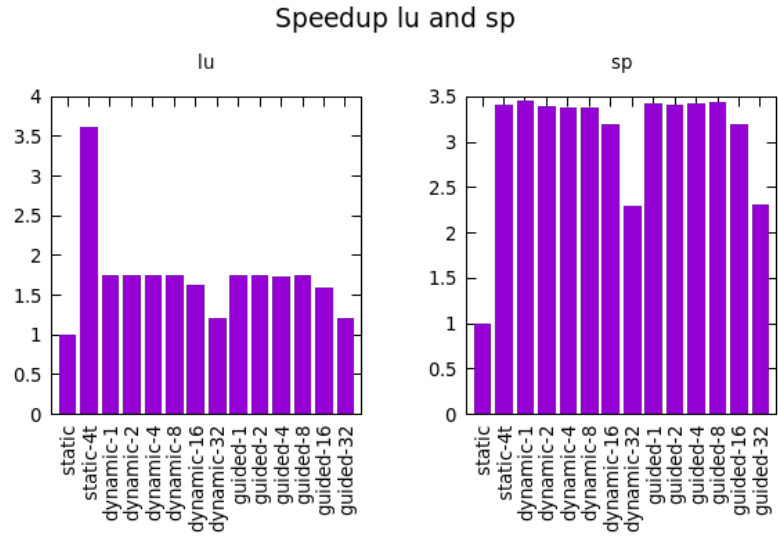


Figure 16: Speedup lu and sp

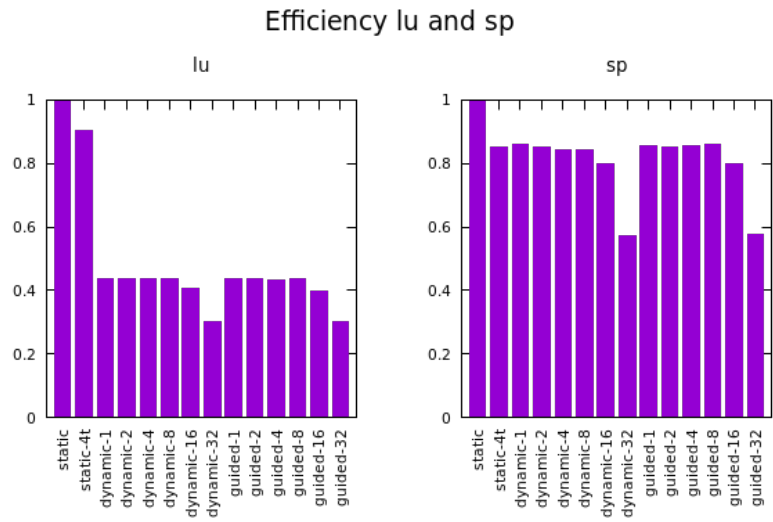


Figure 17: Efficiency lu and sp

LU Kernel

scheduling	num_threads	run_time	speedup	efficiency
static	1	15.477	1	1
static	4	4.292	3.606	0.9015
dynamic_1	4	8.894	1.7401	0.435
dynamic_2	4	8.854	1.748	0.437
dynamic_4	4	8.845	1.7498	0.4374
dynamic_8	4	8.848	1.7492	0.4373
dynamic_16	4	9.51	1.6274	0.4068
dynamic_32	4	12.906	1.1992	0.2998
guided_1	4	8.838	1.7511	0.4377
guided_2	4	8.856	1.7476	0.4369
guided_4	4	8.927	1.7337	0.4334
guided_8	4	8.842	1.7503	0.4375
guided_16	4	9.747	1.5878	0.3969
guided_32	4	12.899	1.1998	0.2999

Table 8: lu scheduling

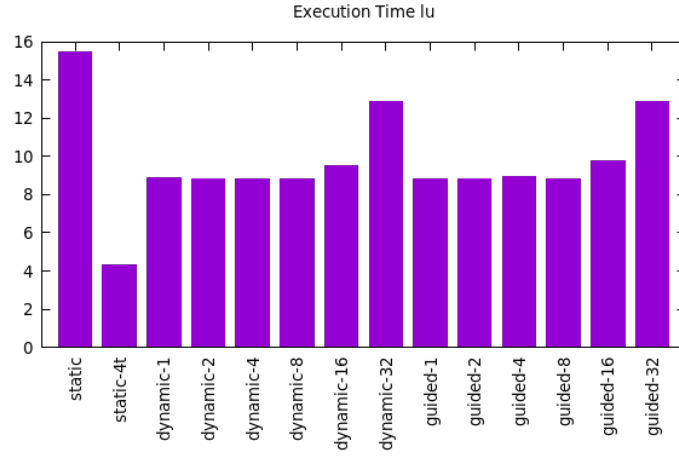


Figure 18: Execution Time lu

SP Kernel

scheduling	num_threads	run_time	speedup	efficiency
static	1	15.661	1	1
static	4	4.596	3.4075	0.8518
dynamic_1	4	4.546	3.445	0.8612
dynamic_2	4	4.611	3.3964	0.8491
dynamic_4	4	4.642	3.3737	0.8434
dynamic_8	4	4.647	3.3701	0.8425
dynamic_16	4	4.906	3.1922	0.798
dynamic_32	4	6.825	2.2946	0.5736
guided_1	4	4.575	3.4231	0.8557
guided_2	4	4.607	3.3993	0.8498
guided_4	4	4.579	3.4201	0.855
guided_8	4	4.558	3.4359	0.8589
guided_16	4	4.911	3.1889	0.7972
guided_32	4	6.813	2.2986	0.5746

Table 9: sp scheduling

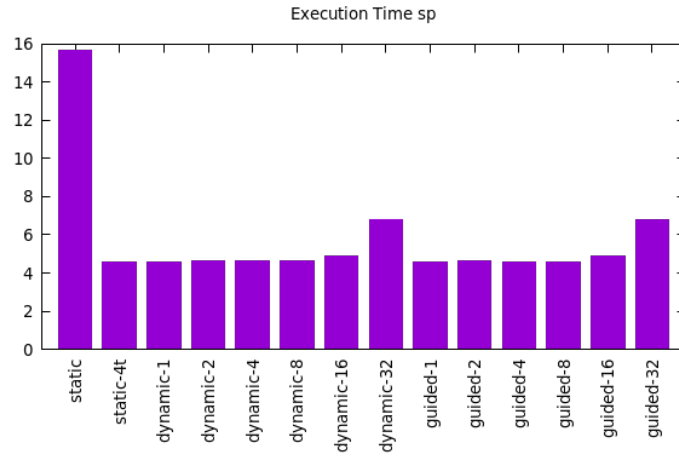


Figure 19: Execution Time sp

1 List of commands

```
# queue jobs
qsub -N hello_omp_v1_3 hello_omp_v1.sge
qsub -N hello_omp_v2_3 -v OMP_NUM_THREADS='3' hello_omp_v2.sge
qsub -N hello_omp_v2_2 -v OMP_NUM_THREADS='2' hello_omp_v2.sge
qsub -N hello_omp_v2_1 -v OMP_NUM_THREADS='1' hello_omp_v2.sge
```

Figure 20: Warm-up Activity

```
# queue jobs
./queue.sh -s sv -p sv_omp -n 1000000000 | sh

# get dat and png files from csv files
./sumarize_sv.sh
```

Figure 21: Question 1

```
# queue jobs
./queue_mm.sh | sh

# get dat and png files from csv files
./sumarize_mm.sh
```

Figure 22: Question 2

```
# queue jobs
./queue_flops.sh -n flops | sh
./queue_flops.sh -n flops2 | sh
```

Figure 23: Question 3

```
# get counter list
qsub papi_avail.sge

# queue jobs
./queue_all_counters.sh | sh

# get csv files
./sumarize_counters.sh
```

Figure 24: Question 5

```
# queue jobs
./queue_npb.sh | sh

# get csv, dat and png files
./sumarize_npb.sh
```

Figure 25: Question 6

```
# get binaries
./build_scheduling.sh

# queue jobs
./queue_scheduling.sh | sh

# csv files
./sumarize_scheduling.sh
```

Figure 26: Question 7